

Simple Exploratory Data Analysis in R

Marius G. Sampid

29 October 2018

Introduction

In this report, I present a simple exploratory data analysis using the publicly available diamond data set, which can be found in <https://vincentarelbundock.github.io/Rdatasets/datasets.html>. The analysis are produced using R markdown script so that readers can easily view the R codes, comments, and output results.

Data Description

carat: weight of the diamond (0.2-5.01),

cut: quality of the cut (Fair, Good, Very Good, Premium, Ideal),

color: diamond colour, from J (worst) to D (best),

clarity: a measurement of how clear the diamond is (I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best)),

depth: total depth percentage = $z / \text{mean}(x, y) = 2 * z / (x + y)$ (43-79),

table: width of top of diamond relative to widest point (43-95)

price: price in US dollars (\$326-\$18,823)

x: length in mm (0-10.74)

y: width in mm (0-58.9)

z: depth in mm (0-31.8)

Reading data into R

```
# Get the address/path where data is stored:  
address <-  
"C:/Users/marius/Desktop/deskTopFiles/DataAnalysis_R_Python/SampleData.csv"  
  
# store data in an object called 'myData'  
myData <- read.csv(address)
```

Exploring the data set

In this section, I present some basic data visualization to see the overall shape and meaning in the data. That is, view the data in some different perspectives.

```
dim(myData) # see the dimension of the data
```

```
## [1] 53940    10
```

```
head(myData) # See the first 6 rows of the data set. A quick snapshot of the data to make sure that the data was correctly imported into R
```

```
##   carat      cut color clarity depth table price     x     y     z
## 1  0.23    Ideal     E    SI2   61.5    55   326  3.95  3.98  2.43
## 2  0.21  Premium     E    SI1   59.8    61   326  3.89  3.84  2.31
## 3  0.23     Good     E    VS1   56.9    65   327  4.05  4.07  2.31
## 4  0.29  Premium     I    VS2   62.4    58   334  4.20  4.23  2.63
## 5  0.31     Good     J    SI2   63.3    58   335  4.34  4.35  2.75
## 6  0.24 Very Good     J   VVS2   62.8    57   336  3.94  3.96  2.48
```

```
str(myData) # see the basic structure of the data
```

```
## 'data.frame':    53940 obs. of  10 variables:
## $ carat : num  0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
## $ cut : Factor w/ 5 levels "Fair","Good",...: 3 4 2 4 2 5 5 5 1 5 ...
## $ color : Factor w/ 7 levels "D","E","F","G",...: 2 2 2 6 7 7 6 5 2 5 ...
## $ clarity: Factor w/ 8 levels "I1","IF","SI1",...: 4 3 5 6 4 8 7 3 6 5 ...
## $ depth : num  61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
## $ table : num  55 61 65 58 58 57 57 55 61 61 ...
## $ price : int  326 326 327 334 335 336 336 337 337 338 ...
## $ x : num  3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
## $ y : num  3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
## $ z : num  2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
```

The output of the `str()` function above shows that the data has 53940 observations and 10 variables. Of the 10 variables, there are 3 factor (categorical) variables of 5, 7, and 8 levels, respectively.

We can easily see the summary statistics of the numeric variables such as the mean, standard deviation, variance, etc with 95% confidence level using the `bsicStats()` function from the `fBasics` package. This function is also important to see if there are any missing values (NA's):

```
install.packages("fBasics") # install package
```

```
library(fBasics) # load package
```

```
# isolate or subset the numeric variables and calculate summary statistics.
```

```
basicStats(myData[,c(1,5:10)])
```

```
##           carat           depth           table           price
## nobs      53940.000000  5.394000e+04  5.394000e+04  5.394000e+04
## NAs        0.000000    0.000000e+00  0.000000e+00  0.000000e+00
## Minimum    0.200000    4.300000e+01  4.300000e+01  3.260000e+02
## Maximum    5.010000    7.900000e+01  9.500000e+01  1.882300e+04
## 1. Quartile 0.400000    6.100000e+01  5.600000e+01  9.500000e+02
## 3. Quartile 1.040000    6.250000e+01  5.900000e+01  5.324250e+03
```

```
## Mean          0.797940  6.174941e+01  5.745718e+01  3.932800e+03
## Median        0.700000  6.180000e+01  5.700000e+01  2.401000e+03
## Sum           43040.870000  3.330763e+06  3.099241e+06  2.121352e+08
## SE Mean       0.002041  6.168000e-03  9.621000e-03  1.717736e+01
## LCL Mean      0.793939  6.173732e+01  5.743833e+01  3.899132e+03
## UCL Mean      0.801940  6.176149e+01  5.747604e+01  3.966467e+03
## Variance      0.224687  2.052404e+00  4.992948e+00  1.591563e+07
## Stdev         0.474011  1.432621e+00  2.234491e+00  3.989440e+03
## Skewness      1.116584 -8.228900e-02  7.968520e-01  1.618305e+00
## Kurtosis      1.256250  5.738447e+00  2.801271e+00  2.177191e+00
##              x              y              z
## nobs          53940.000000  5.394000e+04  5.394000e+04
## NAs           0.000000  0.000000e+00  0.000000e+00
## Minimum       0.000000  0.000000e+00  0.000000e+00
## Maximum       10.740000  5.890000e+01  3.180000e+01
## 1. Quartile   4.710000  4.720000e+00  2.910000e+00
## 3. Quartile   6.540000  6.540000e+00  4.040000e+00
## Mean          5.731157  5.734526e+00  3.538734e+00
## Median        5.700000  5.710000e+00  3.530000e+00
## Sum           309138.620000  3.093203e+05  1.908793e+05
## SE Mean       0.004830  4.918000e-03  3.039000e-03
## LCL Mean      5.721690  5.724887e+00  3.532778e+00
## UCL Mean      5.740624  5.744165e+00  3.544689e+00
## Variance      1.258347  1.304472e+00  4.980110e-01
## Stdev         1.121761  1.142135e+00  7.056990e-01
## Skewness      0.378655  2.434031e+00  1.522338e+00
## Kurtosis      -0.618303  9.120250e+01  4.708029e+01
```

We can also see from the summary statistics that the distribution of the numeric variables: 'carat', 'table', 'price', 'x', 'y' and 'z' are positively skewed while 'depth' is slightly negatively skewed. Apart from variable 'x', the variables also show significant excess kurtosis. Thus the variables are not normally distributed. A normal distribution has excess kurtosis of 0, is symmetric around the mean with 0 skewness. Assuming normality with this data set to predict or forecast future diamond prices might yield false results because extreme events will be neglected or cut out. Therefore, we need to rely on the standardized residuals for forecasting. We also see from the summary statistics that there are no missing values.

It is also worth noting that a positive kurtosis imply that there are more data points in the tails than the normal distribution and a negative kurtosis imply that there are less data points in the tails than the normal distribution, which is the case for variable 'x'.

We can also use the summary function; `summary()`, to see summary statistics of the numeric variables as:

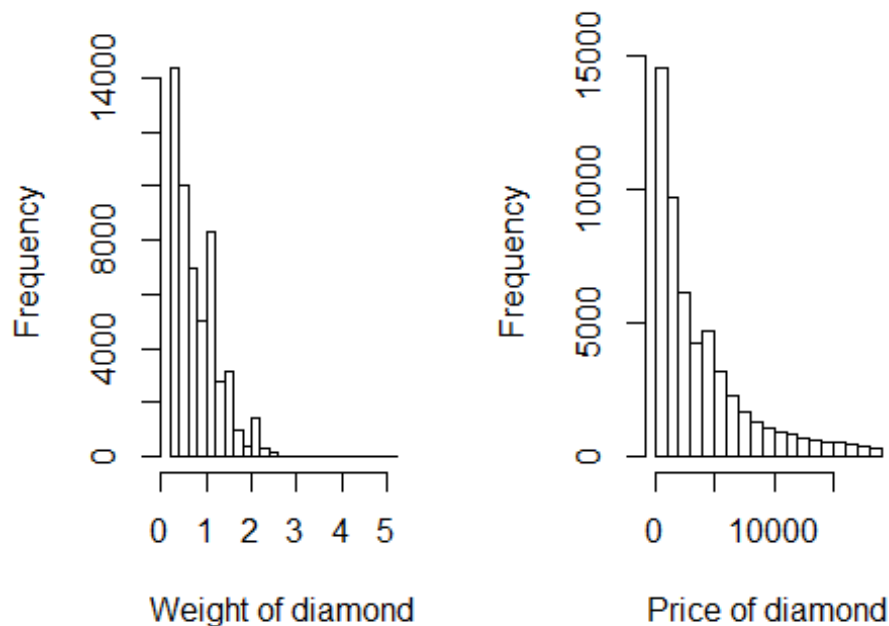
```
summary(myData[,c(1,5:10)]) # give summary of the numeric variables in the data set.
```

```
##      carat      depth      table      price
## Min.   :0.2000   Min.   :43.00   Min.   :43.00   Min.    : 326
## 1st Qu.:0.4000   1st Qu.:61.00   1st Qu.:56.00   1st Qu.: 950
```

```
## Median :0.7000    Median :61.80    Median :57.00    Median : 2401
## Mean   :0.7979    Mean   :61.75    Mean   :57.46    Mean   : 3933
## 3rd Qu.:1.0400    3rd Qu.:62.50    3rd Qu.:59.00    3rd Qu.: 5324
## Max.   :5.0100    Max.   :79.00    Max.   :95.00    Max.   :18823
##      x              y              z
## Min.   : 0.000    Min.   : 0.000    Min.   : 0.000
## 1st Qu.: 4.710    1st Qu.: 4.720    1st Qu.: 2.910
## Median : 5.700    Median : 5.710    Median : 3.530
## Mean   : 5.731    Mean   : 5.735    Mean   : 3.539
## 3rd Qu.: 6.540    3rd Qu.: 6.540    3rd Qu.: 4.040
## Max.   :10.740    Max.   :58.900    Max.   :31.800
```

The distribution of each variable can also be viewed with the `hist()` command. For example, let's take a look at variables 'carat' and 'price'; the weights and price of diamond, respectively:

```
par(mfrow = c(1,2)) # create a matrix of one rows and 2 columns to display
figures
diamondWeight = myData[,1]
diamondPrice  = myData[,7]
hist(diamondWeight, breaks = 20, main = "", xlab = "Weight of diamond")
hist(diamondPrice, breaks = 20, main = "", xlab = "Price of diamond")
```




```
par(mfrow = c(1,1)) # bring back to its original layout
```

We can see from the above plots that the distributions of the weights and prices of diamond are far from a normal distribution. The rest of the plots are shown below:

```

par(mfrow = c(2,4)) # create a matrix of two rows and 4 columns to display
figures
hist(myData[,1], main = "", xlab = "carat", breaks = 20)
hist(myData[,5], main = "", xlab = "depth", breaks = 20)
hist(myData[,6], main = "", xlab = "table", breaks = 20)
hist(myData[,7], main = "", xlab = "price", breaks = 20)
hist(myData[,8], main = "", xlab = "x", breaks = 20)
hist(myData[,9], main = "", xlab = "y", breaks = 20)
hist(myData[,10], main = "", xlab = "z", breaks = 20)
par(mfrow = c(1,1)) # bring back to its original layout

```



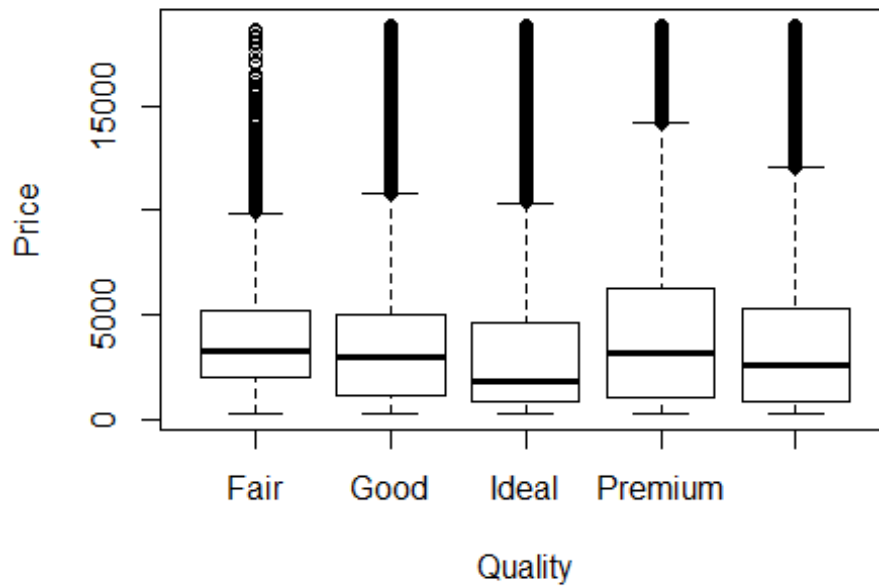
Another good method to visualize the distribution of the data set across the various categories is by using boxplots as follows:

```

boxplot(myData$price~myData$cut, xlab = "Quality", ylab = "Price", main =
"Boxplot distribution of the quality of diamond")

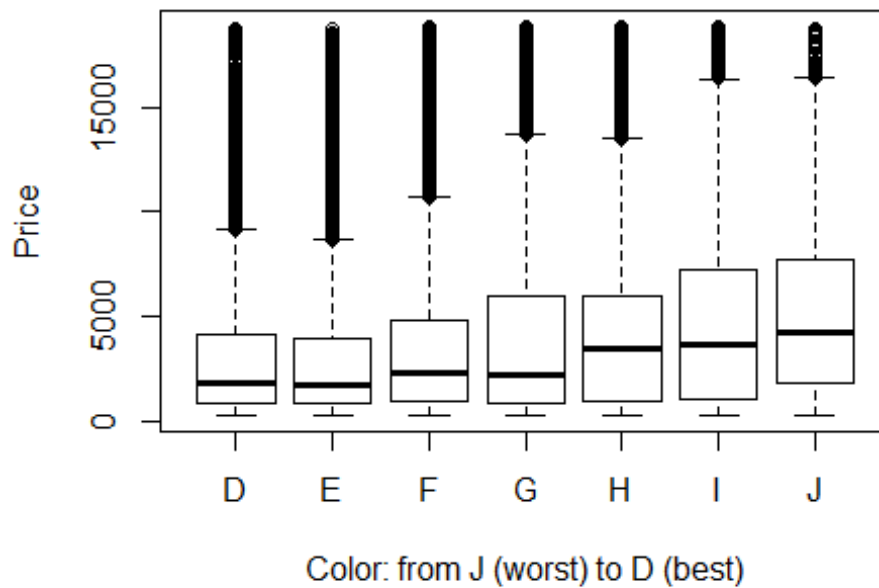
```

Boxplot distribution of the quality of diamond




```
boxplot(myData$price~myData$color, xlab = "Color: from J (worst) to D (best)",  
ylab = "Price", main = "Boxplot distribution of diamond color")
```

Boxplot distribution of diamond color



```
boxplot(myData$price~myData$clarity, xlab = "Clarity: from I1 (worst), SI2, SI1, SV2, VS1, VVS2, VVS1, IF (best)", ylab = "Price", main = "Boxplot distribution of the clarity of diamond")
```



We can see from the boxplots that the distribution of the diamond data set with respect to price in the various categories are all skewed with huge outliers.

The variable names in the data set can be viewed with the `name()` function as seen below. This information will help in case we want to subset the data using `[]` or `$` sign, so that we can easily refer to the variable names.

```
names(myData)

## [1] "carat" "cut" "color" "clarity" "depth" "table" "price"
## [8] "x" "y" "z"
```

We now take a closer look at the factor (Categorical) variables

Of the 53940 observations, how many of the diamonds were recorded as Fair, Good, Very Good, Premium, and Ideal quality? The following code gives us the answer:

```
attach(myData)
table(cut)

## cut
## Fair Good Ideal Premium Very Good
## 1610 4906 21551 13791 12082
```

Note: To avoid using '\$' sign or '[,]' in subsetting the data set, We can use the attach() function to attach everything hidden within the data set to the work space. This makes it easy to work directly with the variable names. Remember to detach at the end using the detach() function.

What was the distribution of the colors from worst (J) to best (D), and the measurement of how clear the diamonds were (I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best))? The following code gives us this answer:

```
table(color) # distribution of diamond color from worst (J) to best (D)

## color
##      D      E      F      G      H      I      J
## 6775  9797  9542 11292  8304  5422  2808

table(clarity) # distribution of clarity of diamond; I1 (worst), SI2, SI1,
VS2, VS1, VVS2, VVS1, IF (best)

## clarity
##      I1      IF     SI1     SI2     VS1     VS2     VVS1     VVS2
##   741   1790 13065   9194   8171 12258   3655   5066
```

To see the total price (in US dollars) on each category of the quality of diamond, diamond color and diamond clarity, we use the following codes:

```
# Calculates the total price (in US dollars) on each category of the quality
of diamond (Fair, Good, Very Good, Premium, Ideal).
TotalPrice_cut <- aggregate(price,by=list(cut),sum)
colnames(TotalPrice_cut) <- c("Quality of Diamond ", " Total Price")
TotalPrice_cut

##   Quality of Diamond   Total Price
## 1                Fair    7017600
## 2                 Good   19275009
## 3                 Ideal   74513487
## 4                Premium   63221498
## 5             Very Good   48107623

# Calculates total price (in US dollars) on each category of the color of
diamond from J (worst) to D (best).
TotalPrice_color <- aggregate(price,by=list(color),sum)
colnames(TotalPrice_color) <- c("Diamond Color ", " Total Pricet")
TotalPrice_color

##   Diamond Color   Total Pricet
## 1              D    21476439
## 2              E    30142944
## 3              F    35542866
## 4              G    45158240
## 5              H    37257301
## 6              I    27608146
## 7              J    14949281
```


Calculates the total price (in US dollars) on each category of the clarity of diamond (Fair, Good, Very Good, Premium, Ideal).

```
TotalPrice_clarity <- aggregate(price,by=list(clarity),sum)
colnames(TotalPrice_clarity) <- c("Clarity of Diamond ", " Total Price")
TotalPrice_clarity
```

##	Clarity of Diamond	Total Price
## 1	I1	2907809
## 2	IF	5128062
## 3	SI1	52207755
## 4	SI2	46549485
## 5	VS1	31372190
## 6	VS2	48112520
## 7	VVS1	9221984
## 8	VVS2	16635412

To see the total weights, in carats, on each category based on the color of diamond, we employ the following code:

Calculates total weight, in carats, of diamond on each category of the color of diamond from J (worst) to D (best).

```
TotalWeight_color <- aggregate(carat,by=list(color),sum)
colnames(TotalWeight_color) <- c("Diamond Color ", " Total Weight")
TotalWeight_color
```

##	Diamond Color	Total Weight
## 1	D	4456.56
## 2	E	6445.12
## 3	F	7028.05
## 4	G	8708.28
## 5	H	7571.58
## 6	I	5568.00
## 7	J	3263.28

To see the total weights, in carats, of each category based on the quality of diamond, we use the following code:

TotalWeightCarat <- aggregate(carat,by=list(cut),sum) # calculates total weight in each category of quality of diamond.

```
colnames(TotalWeightCarat) <- c("Quality ", " Total weight of diamond (carats)")
TotalWeightCarat
```

##	Quality	Total weight of diamond (carats)
## 1	Fair	1684.28
## 2	Good	4166.10
## 3	Ideal	15146.84
## 4	Premium	12300.95
## 5	Very Good	9742.70

Note that we can also use other aggregate functions to get information such as summary statistics, max, mean, median, and counts values, etc. For example, the following codes show summary statistics of price based on quality of diamond, color of diamond and clarity of diamond:

```
aggregate(price,by=list(cut),summary) # summary statistics for the quality of diamond based on price
```

```
##      Group.1 x.Min.  x.1st Qu.  x.Median x.Mean  x.3rd Qu.  x.Max.
## 1      Fair   337      2050      3282    4359      5206    18570
## 2      Good   327      1145      3050    3929      5028    18790
## 3     Ideal   326       878      1810    3458      4678    18810
## 4   Premium   326      1046      3185    4584      6296    18820
## 5 Very Good   336       912      2648    3982      5373    18820
```

```
aggregate(price,by=list(color),summary) # summary statistics for the color of diamond based on price
```

```
##      Group.1 x.Min.  x.1st Qu.  x.Median x.Mean  x.3rd Qu.  x.Max.
## 1          D   357       911      1838    3170      4214    18690
## 2          E   326       882      1739    3077      4003    18730
## 3          F   342       982      2344    3725      4868    18790
## 4          G   354       931      2242    3999      6048    18820
## 5          H   337       984      3460    4487      5980    18800
## 6          I   334      1120      3730    5092      7202    18820
## 7          J   335      1860      4234    5324      7695    18710
```

```
aggregate(price,by=list(clarity),summary) # summary statistics for the clarity of diamond based on price
```

```
##      Group.1  x.Min.  x.1st Qu.  x.Median  x.Mean  x.3rd Qu.  x.Max.
## 1         I1   345.0    2080.0   3344.0   3924.0    5161.0  18530.0
## 2         IF   369.0     895.0   1080.0   2865.0    2388.0  18810.0
## 3        SI1   326.0   1089.0   2822.0   3996.0    5250.0  18820.0
## 4        SI2   326.0   2264.0   4072.0   5063.0    5777.0  18800.0
## 5        VS1   327.0     876.0   2005.0   3839.0    6023.0  18800.0
## 6        VS2   334.0     900.0   2054.0   3925.0    6024.0  18820.0
## 7       VVS1   336.0     816.0   1093.0   2523.0    2379.0  18780.0
## 8       VVS2   336.0     794.2   1311.0   3284.0    3638.0  18770.0
```

A simple t-test

Is the average price per diamond different depending on the quality of the diamond? This question can be answered by conducting a simple t-tests. Let's conduct two simple t-tests; for the first test, we use qualities: "Fair" and "Good". For the second test we used qualities "Fair" and "Ideal" as follows:

```
t.test(price[cut=="Fair"], price[cut=="Good"]) # first t-test
```

```
##
## Welch Two Sample t-test
##
## data: price[cut == "Fair"] and price[cut == "Good"]
## t = 4.1684, df = 2822.3, p-value = 3.16e-05
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 227.6710 632.1156
## sample estimates:
## mean of x mean of y
## 4358.758 3928.864

t.test(price[cut=="Fair"], price[cut=="Ideal"]) # second t-test

##
## Welch Two Sample t-test
##
## data: price[cut == "Fair"] and price[cut == "Ideal"]
## t = 9.7484, df = 1894.8, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 719.9065 1082.5251
## sample estimates:
## mean of x mean of y
## 4358.758 3457.542

detach(myData)
```

Looking at the output of the t-test, for example the first t-test, when the diamond quality is 'Fair', the average price is 4358.758 and when the diamond quality is 'Good', the average price is 3928.864. The p-value for both t-tests suggest that there is significant difference. Therefore, at 95% confidence level, we should reject the null hypothesis and conclude that the average price per diamond is dependent on the quality of the diamond.

Note: p-value is the smallest level of significance in which we will reject the null hypothesis in favour of the alternative hypothesis. At 95% confidence, we should reject the null hypothesis if p-value < 5% significance level.